

Machine Learning Attack on Arbiter PUF

Justin Schubeck, Alex Liu, Yash Bhat

EEE5716 - Introduction to Hardware Security and Trust: Fall 2022

Import necessary packages.

```
In [1]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
from sklearn.model_selection import train_test_split
from sklearn.model_selection import StratifiedKFold
from sklearn.model_selection import GridSearchCV
from sklearn.svm import SVC
from sklearn.linear_model import LogisticRegression
from sklearn.metrics import accuracy_score
from sklearn.model_selection import cross_val_score
from scipy import stats
```

Read in the data, and format to be [-1, 1] instead of [0, 1].

```
In [2]: df = pd.read_excel('CRPSets.xls', header=None)
df = df.replace(0, -1)
df
```

```
Out[2]:
```

	0	1	2	3	4	5	6	7	8	9	...	55	56	57	58	59	60	61	62	63	64
0	-1	-1	1	1	1	-1	1	-1	-1	1	...	1	1	-1	-1	-1	1	-1	1	1	1
1	1	-1	1	-1	1	-1	-1	-1	1	-1	...	1	1	1	-1	-1	-1	-1	-1	1	-1
2	1	1	-1	1	-1	1	1	-1	1	-1	...	1	-1	1	1	1	-1	-1	1	-1	-1
3	-1	1	-1	-1	-1	-1	1	1	-1	1	...	1	1	1	1	1	-1	-1	1	1	-1
4	1	1	-1	-1	-1	1	1	1	-1	1	...	-1	1	-1	1	1	1	1	-1	-1	1
...
11995	-1	1	-1	1	1	1	-1	-1	1	-1	...	1	-1	-1	-1	-1	1	1	-1	-1	-1
11996	-1	1	-1	-1	1	1	-1	1	-1	1	...	1	1	1	-1	-1	1	-1	1	-1	1
11997	1	1	1	-1	-1	-1	1	1	-1	-1	...	1	1	-1	1	-1	1	1	-1	-1	-1
11998	-1	1	-1	1	1	1	1	1	1	1	...	-1	-1	-1	-1	1	-1	1	-1	-1	1
11999	-1	-1	1	1	1	-1	-1	1	1	1	...	1	1	1	-1	-1	1	1	1	-1	-1

12000 rows × 65 columns

Separate the challenges from responses, and extract features.

```
In [3]: challenges = df.to_numpy()[ :, :64]
responses = df.to_numpy()[ :, 64]
features = np.cumprod(np.fliplr(challenges), axis=1, dtype=np.int8)
```

Support Vector Machine

```
In [4]: # Initialize lists for data collection.
best_estimators = []
best_params = []
train_acc = []
test_acc = []
lower_conf = []
upper_conf = []

# Train-Test splits of CRPs.
ratios = np.concatenate((np.array(list(range(25,100,25)))/12000,
                           np.array(list(range(100,1000,100)))/12000,
                           np.array(list(range(1000,12000,1000)))/12000))

# Iterate training models over all splits.
for ratio in ratios:

    # Split the data into training and test splits.
    X_train, X_test, t_train, t_test = train_test_split(features,
                                                         responses,
                                                         train_size=ratio,
                                                         shuffle=True,
                                                         random_state=7,
                                                         stratify=responses)

    # Define a set of hyperparameters to tune.
    param_grid = {'C':np.logspace(-3,3,7),
                  'kernel':['linear','rbf'],
                  'gamma':list(np.logspace(-3,3,7))+['scale','auto']}

    # Define cross validation strategy.
    cv = StratifiedKFold(n_splits=5,
                         shuffle=True,
                         random_state=7)

    # Create hyperparameter testing structure.
    grid_search = GridSearchCV(SVC(),
                               param_grid=param_grid,
                               cv=cv,
                               scoring='accuracy',
                               refit=True,
                               n_jobs=-1)

    # Tune the SVM model.
    grid_search.fit(X_train, t_train)
    best_estimators += [grid_search.best_estimator_]
    best_params += [grid_search.best_params_]

    # Predict training responses.
    y_train = grid_search.best_estimator_.predict(X_train)
    train_acc += [accuracy_score(t_train, y_train)]
```

```

# Predict test responses.
y_test = grid_search.best_estimator_.predict(X_test)
test_acc += [accuracy_score(t_test, y_test)]

# Calculate 10 accuracy scores using cross validation on best model.
scores = cross_val_score(grid_search.best_estimator_,
                          X_train,
                          t_train,
                          scoring='accuracy',
                          cv=10)

# Generate a 95% confidence interval on the test accuracy prediction.
interval = stats.t.interval(0.95,
                             len(scores) - 1,
                             loc=scores.mean(),
                             scale=scores.std(ddof=1)/np.sqrt(len(scores)))
lower_conf += [interval[0]]
upper_conf += [interval[1]]

```

Logging values for each model.

In [19]: `ratios*12000`

Out[19]: `array([25., 50., 75., 100., 200., 300., 400., 500.,
 600., 700., 800., 900., 1000., 2000., 3000., 4000.,
 5000., 6000., 7000., 8000., 9000., 10000., 11000.])`

In [5]: `best_params`

Out[5]: `[{'C': 0.1, 'gamma': 0.001, 'kernel': 'linear'},
{'C': 1.0, 'gamma': 0.01, 'kernel': 'rbf'},
{'C': 1.0, 'gamma': 0.01, 'kernel': 'rbf'},
{'C': 10.0, 'gamma': 0.001, 'kernel': 'rbf'},
{'C': 10.0, 'gamma': 0.01, 'kernel': 'rbf'},
{'C': 10.0, 'gamma': 0.01, 'kernel': 'rbf'},
{'C': 0.1, 'gamma': 0.001, 'kernel': 'linear'},
{'C': 1000.0, 'gamma': 0.001, 'kernel': 'rbf'},
{'C': 0.1, 'gamma': 0.001, 'kernel': 'linear'},
{'C': 10.0, 'gamma': 0.001, 'kernel': 'linear'},
{'C': 1000.0, 'gamma': 0.001, 'kernel': 'rbf'},
{'C': 10.0, 'gamma': 0.001, 'kernel': 'linear'},
{'C': 0.1, 'gamma': 0.001, 'kernel': 'linear'},
{'C': 10.0, 'gamma': 0.001, 'kernel': 'linear'},
{'C': 100.0, 'gamma': 0.001, 'kernel': 'linear'},
{'C': 100.0, 'gamma': 0.001, 'kernel': 'linear'},
{'C': 100.0, 'gamma': 0.001, 'kernel': 'linear'},
{'C': 100.0, 'gamma': 0.001, 'kernel': 'linear'},
{'C': 1000.0, 'gamma': 0.001, 'kernel': 'linear'},
{'C': 1000.0, 'gamma': 0.001, 'kernel': 'linear'},
{'C': 100.0, 'gamma': 0.001, 'kernel': 'linear'},
{'C': 100.0, 'gamma': 0.001, 'kernel': 'linear'},
{'C': 1000.0, 'gamma': 0.001, 'kernel': 'linear'}]`

In [6]: `best_estimators`

```
Out[6]: [SVC(C=0.1, gamma=0.001, kernel='linear'),
SVC(gamma=0.01),
SVC(gamma=0.01),
SVC(C=10.0, gamma=0.001),
SVC(C=10.0, gamma=0.01),
SVC(C=10.0, gamma=0.01),
SVC(C=0.1, gamma=0.001, kernel='linear'),
SVC(C=1000.0, gamma=0.001),
SVC(C=0.1, gamma=0.001, kernel='linear'),
SVC(C=10.0, gamma=0.001, kernel='linear'),
SVC(C=1000.0, gamma=0.001),
SVC(C=10.0, gamma=0.001, kernel='linear'),
SVC(C=0.1, gamma=0.001, kernel='linear'),
SVC(C=10.0, gamma=0.001, kernel='linear'),
SVC(C=100.0, gamma=0.001, kernel='linear'),
SVC(C=100.0, gamma=0.001, kernel='linear'),
SVC(C=100.0, gamma=0.001, kernel='linear'),
SVC(C=100.0, gamma=0.001, kernel='linear'),
SVC(C=1000.0, gamma=0.001, kernel='linear'),
SVC(C=1000.0, gamma=0.001, kernel='linear'),
SVC(C=100.0, gamma=0.001, kernel='linear'),
SVC(C=100.0, gamma=0.001, kernel='linear'),
SVC(C=1000.0, gamma=0.001, kernel='linear')]
```

```
In [7]: train_acc
```

```
Out[7]: [1.0,
1.0,
1.0,
0.98,
1.0,
1.0,
0.98,
1.0,
0.9816666666666667,
1.0,
1.0,
1.0,
0.974,
0.9985,
1.0,
1.0,
1.0,
1.0,
1.0,
1.0,
0.9995555555555555,
0.9992,
1.0]
```

```
In [8]: test_acc
```

```
Out[8]: [0.6982881002087683,  
0.6758995815899581,  
0.7477568134171908,  
0.8163865546218487,  
0.8680508474576272,  
0.8804273504273504,  
0.9043103448275862,  
0.935304347826087,  
0.9363157894736842,  
0.9567256637168141,  
0.9558035714285714,  
0.9548648648648649,  
0.9532727272727273,  
0.988,  
0.9885555555555555,  
0.991875,  
0.9941428571428571,  
0.9945,  
0.9938,  
0.995,  
0.9946666666666667,  
0.9965,  
0.998]
```

```
In [9]: lower_conf
```

```
Out[9]: [0.5454152743753193,  
0.559351617987147,  
0.6538312283302322,  
0.6374809888837991,  
0.8817859431399552,  
0.8554420710059591,  
0.8886818522171743,  
0.9094108474732326,  
0.9186621982690928,  
0.9411651366401025,  
0.9414965165242757,  
0.9482579420118066,  
0.945790205367151,  
0.9776570108525163,  
0.9850089307290959,  
0.9923340222168601,  
0.9933265242603575,  
0.9940599631074988,  
0.994015875752269,  
0.9952185123610476,  
0.9963813497567361,  
0.9956122023325912,  
0.9957267769953407]
```

```
In [10]: upper_conf
```

```
Out[10]: [0.9545847256246807,
0.8006483820128528,
0.8211687716697679,
0.8825190111162009,
0.9382140568600447,
0.9378912623273745,
0.9513181477828255,
0.9425891525267672,
0.9513378017309071,
0.9731205776456117,
0.9710034834757244,
0.9695198357659712,
0.9682097946328487,
0.9873429891474835,
0.9936577359375707,
0.9976659777831397,
0.9982734757396425,
0.9982733702258345,
0.9985555528191594,
0.9982814876389526,
0.9985075391321526,
0.9989877976674085,
0.9984550411864778]
```


Visual display of SVM results.

```
In [15]: plt.figure(figsize=(10,6));

# Plot training accuracies.
plt.plot(ratios*12000, train_acc, color='red');
plt.scatter(ratios*12000, train_acc, color='red', label='Training');

# Plot test accuracies.
plt.plot(ratios*12000, test_acc, color='blue');
plt.scatter(ratios*12000, test_acc, color='blue', label='Test');

plt.title('Training Set Size vs. Accuracy', size=18);
plt.xlabel('Number of CRPs in Training', size=15);
plt.ylabel('Accuracy', size=15);
plt.legend(prop={'size': 15});
```



Logistic Regression

```
In [27]: # Initialize lists for data collection.
best_estimators = []
best_params = []
train_acc = []
test_acc = []
lower_conf = []
upper_conf = []

# Train-Test splits of CRPs.
ratios = np.concatenate((np.array(list(range(25,100,25)))/12000,
                           np.array(list(range(100,1000,100)))/12000,
                           np.array(list(range(1000,12000,1000)))/12000))

# Iterate training models over all splits.
for ratio in ratios:

    # Split the data into training and test splits.
    X_train, X_test, t_train, t_test = train_test_split(features,
                                                         responses,
                                                         train_size=ratio,
                                                         shuffle=True,
                                                         random_state=7,
                                                         stratify=responses)

    # Define a set of hyperparameters to tune.
    param_grid = {'penalty':['l2', 'none'],
                  'C':np.logspace(-3,3,7),
                  'solver':['newton-cg', 'lbfgs', 'sag', 'saga'],
                  'max_iter':[3000]}

    # Define cross validation strategy.
```

```

cv = StratifiedKFold(n_splits=5,
                     shuffle=True,
                     random_state=7)

# Create hyperparameter testing structure.
import warnings
warnings.filterwarnings("ignore", message="Setting penalty='none' will ignore the
grid_search = GridSearchCV(LogisticRegression(),
                           param_grid=param_grid,
                           cv=cv,
                           scoring='accuracy',
                           refit=True,
                           n_jobs=-1)

# Tune the LR model.
grid_search.fit(X_train, t_train)
best_estimators += [grid_search.best_estimator_]
best_params += [grid_search.best_params_]

# Predict training responses.
y_train = grid_search.best_estimator_.predict(X_train)
train_acc += [accuracy_score(t_train, y_train)]

# Predict test responses.
y_test = grid_search.best_estimator_.predict(X_test)
test_acc += [accuracy_score(t_test, y_test)]

# Calculate 10 accuracy scores using cross validation on best model.
scores = cross_val_score(grid_search.best_estimator_,
                          X_train,
                          t_train,
                          scoring='accuracy',
                          cv=10)

# Generate a 95% confidence interval on the test accuracy prediction.
interval = stats.t.interval(0.95,
                             len(scores) - 1,
                             loc=scores.mean(),
                             scale=scores.std(ddof=1)/np.sqrt(len(scores)))
lower_conf += [interval[0]]
upper_conf += [interval[1]]

```

Logging values for each model.

In [28]: ratios*12000

Out[28]: array([25., 50., 75., 100., 200., 300., 400., 500.,
600., 700., 800., 900., 1000., 2000., 3000., 4000.,
5000., 6000., 7000., 8000., 9000., 10000., 11000.])

In [29]: best_params


```
Out[29]: [{'C': 0.001, 'max_iter': 3000, 'penalty': 'none', 'solver': 'newton-cg'},
{'C': 0.001, 'max_iter': 3000, 'penalty': 'none', 'solver': 'newton-cg'},
{'C': 1.0, 'max_iter': 3000, 'penalty': 'l2', 'solver': 'newton-cg'},
{'C': 1.0, 'max_iter': 3000, 'penalty': 'none', 'solver': 'sag'},
{'C': 0.001, 'max_iter': 3000, 'penalty': 'none', 'solver': 'newton-cg'},
{'C': 1.0, 'max_iter': 3000, 'penalty': 'l2', 'solver': 'newton-cg'},
{'C': 0.001, 'max_iter': 3000, 'penalty': 'none', 'solver': 'saga'},
{'C': 0.001, 'max_iter': 3000, 'penalty': 'none', 'solver': 'newton-cg'},
{'C': 10.0, 'max_iter': 3000, 'penalty': 'l2', 'solver': 'newton-cg'},
{'C': 0.001, 'max_iter': 3000, 'penalty': 'none', 'solver': 'saga'},
{'C': 0.001, 'max_iter': 3000, 'penalty': 'none', 'solver': 'lbfgs'},
{'C': 10.0, 'max_iter': 3000, 'penalty': 'l2', 'solver': 'newton-cg'},
{'C': 0.001, 'max_iter': 3000, 'penalty': 'none', 'solver': 'lbfgs'},
{'C': 100.0, 'max_iter': 3000, 'penalty': 'l2', 'solver': 'newton-cg'},
{'C': 100.0, 'max_iter': 3000, 'penalty': 'l2', 'solver': 'newton-cg'},
{'C': 0.001, 'max_iter': 3000, 'penalty': 'none', 'solver': 'sag'},
{'C': 0.001, 'max_iter': 3000, 'penalty': 'none', 'solver': 'newton-cg'},
{'C': 0.001, 'max_iter': 3000, 'penalty': 'none', 'solver': 'newton-cg'},
{'C': 1000.0, 'max_iter': 3000, 'penalty': 'l2', 'solver': 'newton-cg'},
{'C': 1000.0, 'max_iter': 3000, 'penalty': 'l2', 'solver': 'newton-cg'},
{'C': 0.001, 'max_iter': 3000, 'penalty': 'none', 'solver': 'saga'},
{'C': 0.001, 'max_iter': 3000, 'penalty': 'none', 'solver': 'sag'},
{'C': 0.001, 'max_iter': 3000, 'penalty': 'none', 'solver': 'lbfgs'}]
```

```
In [30]: best_estimators
```

```
Out[30]: [LogisticRegression(C=0.001, max_iter=3000, penalty='none', solver='newton-cg'),
LogisticRegression(C=0.001, max_iter=3000, penalty='none', solver='newton-cg'),
LogisticRegression(max_iter=3000, solver='newton-cg'),
LogisticRegression(max_iter=3000, penalty='none', solver='sag'),
LogisticRegression(C=0.001, max_iter=3000, penalty='none', solver='newton-cg'),
LogisticRegression(max_iter=3000, solver='newton-cg'),
LogisticRegression(C=0.001, max_iter=3000, penalty='none', solver='saga'),
LogisticRegression(C=0.001, max_iter=3000, penalty='none', solver='newton-cg'),
LogisticRegression(C=10.0, max_iter=3000, solver='newton-cg'),
LogisticRegression(C=0.001, max_iter=3000, penalty='none', solver='saga'),
LogisticRegression(C=0.001, max_iter=3000, penalty='none'),
LogisticRegression(C=10.0, max_iter=3000, solver='newton-cg'),
LogisticRegression(C=0.001, max_iter=3000, penalty='none'),
LogisticRegression(C=100.0, max_iter=3000, solver='newton-cg'),
LogisticRegression(C=100.0, max_iter=3000, solver='newton-cg'),
LogisticRegression(C=0.001, max_iter=3000, penalty='none', solver='sag'),
LogisticRegression(C=0.001, max_iter=3000, penalty='none', solver='newton-cg'),
LogisticRegression(C=0.001, max_iter=3000, penalty='none', solver='newton-cg'),
LogisticRegression(C=1000.0, max_iter=3000, solver='newton-cg'),
LogisticRegression(C=1000.0, max_iter=3000, solver='newton-cg'),
LogisticRegression(C=0.001, max_iter=3000, penalty='none', solver='saga'),
LogisticRegression(C=0.001, max_iter=3000, penalty='none', solver='sag'),
LogisticRegression(C=0.001, max_iter=3000, penalty='none')]
```

```
In [31]: train_acc
```

[illegible]

```
In [32]: test_acc
```

```
Out[32]: [0.7064718162839249,
0.717489539748954,
0.7961425576519916,
0.8280672268907563,
0.8630508474576272,
0.8996581196581197,
0.9168103448275862,
0.9224347826086956,
0.9471929824561404,
0.9567256637168141,
0.963125,
0.9635135135135136,
0.9737272727272728,
0.9911,
0.9895555555555555,
0.994,
0.9941428571428571,
0.9933333333333333,
0.995,
0.9955,
0.9943333333333333,
0.9975,
0.997]
```

```
In [33]: lower_conf
```

```
Out[33]: [0.5454152743753193,
0.6708899571190191,
0.657147484132205,
0.690267921872878,
0.8736818522171743,
0.8664505393290493,
0.9032312082271344,
0.9147645803759583,
0.9195801648139663,
0.9467498806607417,
0.9586741295245832,
0.955425917351829,
0.9459351617987146,
0.9781547311656396,
0.9854874466840907,
0.9923340222168601,
0.993488653314775,
0.9945738777623125,
0.9946734824651404,
0.9968689214186295,
0.9961057079276706,
0.9947552544262747,
0.9957538621173503]
```

```
In [34]: upper_conf
```

```
Out[34]: [0.9545847256246807,
0.8091100428809809,
0.8178525158677951,
0.9497320781271219,
0.9363181477828255,
0.9468827940042844,
0.9667687917728655,
0.9652354196240418,
0.9604198351860334,
0.9818215479106871,
0.9688258704754166,
0.9779074159815041,
0.9700648382012851,
0.9888452688343603,
0.9965125533159097,
0.9976659777831397,
0.9977113466852249,
0.998092788904354,
0.9978979461062885,
0.9991310785813705,
0.9981165142945515,
0.9990447455737251,
0.9980643197008314]
```

Visual display of LR results.


```
In [35]: plt.figure(figsize=(10,6));

# Plot training accuracies.
plt.plot(ratios*12000, train_acc, color='red');
plt.scatter(ratios*12000, train_acc, color='red', label='Training');

# Plot test accuracies.
plt.plot(ratios*12000, test_acc, color='blue');
```

```
plt.scatter(ratios*12000, test_acc, color='blue', label='Test');

plt.title('Training Set Size vs. Accuracy', size=18);
plt.xlabel('Number of CRPs in Training', size=15);
plt.ylabel('Accuracy', size=15);
plt.legend(prop={'size': 15});
```



In []: